

Deep Learning Action Recognition

"Drinking-Cooking"

Mehdi Gorjian

1 Topic

This is a Deep Learning Action Recognition project which was trained on a GPU server to recognize "Drinking Activity".

2 Dataset

I have implemented different sources to prepare the dataset such as Kinetics, ActivityNet and Youtube. The files were videos with the approximate length of 10 sec.

I wrote a code to extract the frames from the videos, put them in a specific folder and rename them. I extracated 3907 images for drinking and 7831 images for cooking as the training dataset with Data Augmentation.

3 DNN Model

3.1 Architecture

I got the benefit of utilizing Transfer Learning - Fine Tuning from ResNet50 trained on the "ImageNet". Then applied a CNN to it in order to be trained on my custom dataset. For prediction part, I took advantage of the temporal nature of videos, specifically the assumption that subsequent frames in a video will have similar semantic contents. (I used 128 frames in the sequence.)

```
1 baseModel = ResNet50(include_top=False, weights="imagenet", input_tensor=Input(shape
    =(224, 224, 3)), input_shape=(244,244,3))
2 headModel = baseModel.output
3 headModel = AveragePooling2D(pool_size=(7, 7))(headModel)
4 headModel = Flatten()(headModel)
5 headModel = Dense(512, activation="relu")(headModel)
6 headModel = Dropout(0.5)(headModel)
7 headModel = Dense(len(lb.classes_), activation="softmax")(headModel)
8 model = Model(inputs=baseModel.input, outputs=headModel)
9 for layer in baseModel.layers:
10     layer.trainable = False
```

3.2 Input Tensor Shape

Drinking Tensor Shape:

$X_{Train} = (5320, 224, 224, 3)$

$X_{Test} = (1773, 224, 224, 3)$

Cooking Tensor Shape:

$X_{Train} = (5873, 224, 224, 3)$

$X_{Test} = (1958, 224, 224, 3)$

3.3 Layer Structure

Drinking Layer Structure:

- AveragePooling2D Layer: Pool(7,7)
- Flatten Layer
- Dense Layer: 512 + Regularization l1
- Dropout Layer: 50%
- Softmax Prediction

Cooking Layer Structure:

- AveragePooling2D Layer: Pool(7,7)
- Flatten Layer
- Dense Layer: 512 + Regularization l1
- Dropout Layer: 50%
- Softmax Prediction

4 Hyper-parameters

4.1 List and Range of Hyper-parameters

Batch Size	32, 64 & 128
Epoch	10 - 50
Dropout	0.1 - 0.5

4.2 Optimal Hyper-parameters

Batch Size	32
Epoch	35
Dropout	0.5

4.3 Improvement and Performance

- Keras callbacks can help to fix bugs more quickly, and to build better models. They can help to visualize how the model's training is going, and can even help prevent overfitting by implementing checkpoints or customizing the learning rate on each iteration. A callback is a set of functions to be applied at given stages of the training procedure.
- I used a ModelCheckpoint callback to save the model after each epoch and to check the performance of the model. Before utilizing the checkpoint, I set the epochs to 50, however the checkpoint technique, illustrated that there is no improvement after epoch 10.
- I implemented EarlyStopping callback to adjust the optimized Epochs numbers.
- Adding (l1) Regularizaion to the FC layer. It had a good effects on improvement to the project.
- Adding ModelCheckpoint
- Adding RemoteMonitor
- Adding ReduceLROnPlateau
- Meanwhile, I got the benefits of TensorBoard visualization tool to track the training process in live.
- Reducing prediction flickerring
- Changing prediction to show "computing..." text when prediction is out of the boundary
- Revising the dataset: adding and removing some photos and balance them
- Bug fixing

5 Annotated Code

5.1 Training Code

```
1 # USAGE
2 # TO RUN THE CODE: python train.py --data dataset --model output/activity.model --
   label_bin output/lb.pickle --plot output/fig_v1.png --epochs 50
3 # set the matplotlib backend so figures can be saved in the background
4 import matplotlib
5 matplotlib.use("Agg")
6 # import the necessary packages
7 from keras.preprocessing.image import ImageDataGenerator
8 from keras.layers.pooling import AveragePooling2D
9 from keras.applications import ResNet50
10 from keras.layers.core import Dropout
11 from keras.layers.core import Flatten
12 from keras.layers.core import Dense
13 from keras.layers import Input
14 from keras.models import Model
15 from keras.optimizers import SGD
```

```

16 from keras import regularizers
17 from keras.utils import to_categorical
18 from keras.callbacks import EarlyStopping, TensorBoard, History, ModelCheckpoint,
    RemoteMonitor, ReduceLROnPlateau
19 from sklearn.preprocessing import LabelBinarizer
20 from sklearn.model_selection import train_test_split
21 from sklearn.metrics import classification_report, confusion_matrix
22 from imutils import paths
23 import matplotlib.pyplot as plt
24 import numpy as np
25 import argparse
26 import warnings
27 import pickle
28 import time
29 import cv2
30 import os
31 BATCH_SIZE = 32
32 # construct the argument parser and parse the arguments
33 parse = argparse.ArgumentParser()
34 parse.add_argument("-d", "--data", required=True,
35     help="path to input dataset")
36 parse.add_argument("-m", "--model", required=True,
37     help="path to output serialized model")
38 parse.add_argument("-l", "--label_bin", required=True,
39     help="path to output label binarizer")
40 parse.add_argument("-e", "--epochs", type=int, default=25,
41     help="# of epochs to train our network for")
42 parse.add_argument("-p", "--plot", type=str, default="plot.png",
43     help="path to output loss/accuracy plot")
44 args = parse.parse_args()
45
46 # initialize the set of labels from the spots activity dataset we are
47 # going to train our network on
48 LABELS = {"cooking", "drinking"}
49
50 # grab the list of images in our dataset directory, then initialize
51 # the list of data (i.e., images) and class images
52 print("[INFO] loading images...")
53 imagePath = list(paths.list_images(args.data))
54
55 data = []
56 labels = []
57
58 # loop over the image paths
59 for imagePath in imagePath:
60     # extract the class label from the filename
61     label = imagePath.split(os.path.sep)[-2]
62     info_file_name = imagePath.split(os.path.sep)[-1]
63
64     # if the label of the current image is not part of of the labels
65     # are interested in, then ignore the image
66     if label not in LABELS:
67         continue
68
69     # load the image, convert it to RGB channel ordering, and resize
70     # it to be a fixed 224x224 pixels, ignoring aspect ratio
71     image = cv2.imread(imagePath)
72     image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

```

```

73     image = cv2.resize(image, (224, 224))
74     print('[INFO] processing image: ', info_file_name)
75     # update the data and labels lists, respectively
76     data.append(image)
77     labels.append(label)
78
79 # convert the data and labels to NumPy arrays
80 data = np.array(data)
81 labels = np.array(labels)
82 # print('labels before one-hot: ', labels)
83
84 # perform one-hot encoding on the labels
85 lb = LabelBinarizer()
86 labels = lb.fit_transform(labels)
87 labels = to_categorical(labels)
88 # print('labels after one-hot: ', labels)
89
90 # partition the data into training and testing splits using 75% of
91 # the data for training and the remaining 25% for testing
92 print('[INFO] splitting data...')
93
94 trainX, testX, trainY, testY = train_test_split(data, labels,
95     test_size=0.25, stratify=labels, random_state=42)
96
97 print('[INFO] creating generator object...')
98 # initialize the training data augmentation object
99 trainAug = ImageDataGenerator(
100     rotation_range=30,
101     zoom_range=0.15,
102     width_shift_range=0.2,
103     height_shift_range=0.2,
104     shear_range=0.15,
105     horizontal_flip=True,
106     fill_mode="nearest")
107
108 # initialize the validation/testing data augmentation object (which
109 # we'll be adding mean subtraction to)
110 valAug = ImageDataGenerator()
111
112 # define the ImageNet mean subtraction (in RGB order) and set the
113 # the mean subtraction value for each of the data augmentation
114 # objects
115 mean = np.array([123.68, 116.779, 103.939], dtype="float32")
116 trainAug.mean = mean
117 valAug.mean = mean
118
119 # train and test generators
120 train_generator = trainAug.flow(trainX, trainY, batch_size=BATCH_SIZE)
121 validation_generator = valAug.flow(testX, testY)
122 # load the ResNet-50 network, ensuring the head FC layer sets are left
123 # off
124 # ignore warnings
125 warnings.filterwarnings("ignore")
126
127 # using tensorboard
128 tensorboard = TensorBoard(log_dir=f'logs/{time.time()}', batch_size=BATCH_SIZE)
129
130 # baseModel = ResNet50(include_top=False, weights="imagenet",

```

```

131     # input_tensor=Input(shape=(224, 224, 3)))
132
133 baseModel = ResNet50(include_top=False, weights="imagenet", input_tensor=Input(shape
    =(224, 224, 3)), input_shape=(244,244,3))
134
135 # construct the head of the model that will be placed on top of the
136 # the base model
137 headModel = baseModel.output
138 headModel = AveragePooling2D(pool_size=(7, 7))(headModel)
139 headModel = Flatten()(headModel)
140 headModel = Dense(512, activation="relu", kernel_regularizer=regularizers.l1(0.01))(
    headModel)
141 headModel = Dropout(0.5)(headModel)
142 headModel = Dense(len(lb.classes_), activation="softmax")(headModel)
143
144 # place the head FC model on top of the base model (this will become
145 # the actual model we will train)
146 model = Model(inputs=baseModel.input, outputs=headModel)
147
148 # loop over all layers in the base model and freeze them so they will
149 # *not* be updated during the training process
150 for layer in baseModel.layers:
151     layer.trainable = False
152
153 # compile our model (this needs to be done after our setting our
154 # layers to being non-trainable)
155 print("[INFO] compiling model...")
156 opt = SGD(lr=1e-4, momentum=0.9, decay=1e-4 / args.epochs)
157 model.compile(loss="categorical_crossentropy", optimizer=opt,
158     metrics=["accuracy"])
159
160 reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2, patience=5, min_lr
    =0.001)
161 es_callbacks = EarlyStopping(monitor='val_loss', mode='min', restore_best_weights=
    True, patience=10)
162 check_point = ModelCheckpoint(filepath='weights/weights.hdf5', save_best_only=True,
    monitor='val_loss', mode='min', verbose=1)
163
164 # train the head of the network for a few epochs (all other layers
165 # are frozen) -- this will allow the new FC layers to start to become
166 # initialized with actual "learned" values versus pure random
167 print("[INFO] training head...")
168 H = History()
169 # R = RemoteMonitor(root='http://localhost:9000', path='output_temp/', field='data')
170 model.fit_generator(train_generator, steps_per_epoch=len(trainX) // BATCH_SIZE,
    validation_data=validation_generator, validation_steps=len(testX) // BATCH_SIZE,
    epochs=args.epochs, callbacks=[H, tensorboard, check_point, es_callbacks,
    reduce_lr])
171
172 # evaluate the network
173 print("[INFO] evaluating network...")
174 predictions = model.predict(testX, batch_size=BATCH_SIZE)
175
176 print(classification_report(testY.argmax(axis=1),
    predictions.argmax(axis=1), target_names=lb.classes_))
177
178 # plot the training loss and accuracy
179 plt.style.use("ggplot")

```

```

181 plt.figure()
182 plt.plot(H.history["loss"], label="train_loss")
183 plt.plot(H.history["val_loss"], label="val_loss")
184 plt.plot(H.history["accuracy"], label="train_acc")
185 plt.plot(H.history["val_accuracy"], label="val_acc")
186
187 plt.title("Training Loss and Accuracy on Dataset")
188 plt.xlabel("Epoch #")
189 plt.ylabel("Loss | Accuracy")
190 plt.legend(loc="lower left")
191 plt.savefig(args.plot)
192 # serialize the model to disk
193 print("[INFO] serializing network...")
194 model.save(args.model)
195
196 # serialize the label binarizer to disk
197 f = open(args.label_bin, "wb")
198 f.write(pickle.dumps(lb))
199 f.close()
200 print('[INFO] model saved!')
201
202 # -----TESTING RUNS
203 #RUN ON THE FLOYDHUB: python train.py --data /floyd/input/dlaction --model output/
    activity.model --label_bin output/lb.pickle --plot output/fig_v1.png --epochs 50
204
205 # TO DEBUG ON THE LOCAL SYSTEM: python train.py --data dataset_temp --model
    output_temp/activity.model --label_bin output_temp/lb.pickle --plot output_temp/
    _plot1.png --epochs 1

```

5.2 Prediction Code

```

1 # USAGE
2 # python predict_video.py --model model/activity.model --label-bin model/lb.pickle
    --input example_clips/drink1.mp4 --output output/drink1_128avg.avi --action
    drinking --size 128
3
4 # import the necessary packages
5 from keras.models import load_model
6 from collections import deque
7 from utils.json_export import to_json_file
8 from utils.fig_export import fig_plot
9 import numpy as np
10 import argparse
11 import pickle
12 import json
13 import cv2
14
15 # construct the argument parser and parse the arguments
16 ap = argparse.ArgumentParser()
17 ap.add_argument("-m", "--model", required=True,
    help="path to trained serialized model")
18 ap.add_argument("-l", "--label-bin", required=True,
    help="path to label binarizer")
19 ap.add_argument("-i", "--input", required=True,
    help="path to our input video")
20 ap.add_argument("-o", "--output", required=True,
    help="path to our output video")
21 ap.add_argument("-s", "--size", type=int, default=128,

```

```

26     help="size of queue for averaging")
27 ap.add_argument("-a", "--action", required=True, help="choose a predictive action
    from the list [drinking, cooking]")
28 args = ap.parse_args()
29
30 # load the trained model and label binarizer from disk
31 print("[INFO] loading model and label binarizer...")
32 model = load_model(args.model)
33 lb = pickle.loads(open(args.label_bin, "rb").read())
34
35 # initialize the image mean for mean subtraction along with the
36 # predictions queue
37 mean = np.array([123.68, 116.779, 103.939][::1], dtype="float32")
38 Q = deque(maxlen=args.size)
39
40 # initialize the video stream, pointer to output video file, and
41 # frame dimensions
42 vs = cv2.VideoCapture(args.input)
43 writer = None
44 (W, H) = (None, None)
45
46
47 # list of lists [time(milsec), binary_label]
48 timeStamps = []
49 # loop over frames from the video file stream
50 while True:
51     # read the next frame from the file
52     (grabbed, frame) = vs.read()
53
54     # if the frame was not grabbed, then we have reached the end
55     # of the stream
56     if not grabbed:
57         break
58
59     # if the frame dimensions are empty, grab them
60     if W is None or H is None:
61         (H, W) = frame.shape[:2]
62
63     # clone the output frame, then convert it from BGR to RGB
64     # ordering, resize the frame to a fixed 224x224, and then
65     # perform mean subtraction
66     output = frame.copy()
67     # output = cv2.resize(frame.copy(), (244,244))
68     frame = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
69     frame = cv2.resize(frame, (224, 224)).astype("float32")
70     frame -= mean
71
72     # make predictions on the frame and then update the predictions
73     # queue
74     preds = model.predict(np.expand_dims(frame, axis=0))[0]
75     Q.append(preds)
76
77     # perform prediction averaging over the current history of
78     # previous predictions
79     results = np.array(Q).mean(axis=0)
80     i = np.argmax(results)
81     label = lb.classes_[i]
82

```



```

83     if label == args.action:
84         # get the frame time
85         ts = round(vs.get(cv2.CAP_PROP_POS_MSEC), 15)
86         # appending time, label to the list
87         timeStamps.append([round(ts, 16), round(max(results), 5)])
88
89     # draw the activity on the output frame
90     text = "activity: {}".format(label)
91     cv2.putText(output, text, (10, 25), cv2.FONT_HERSHEY_SIMPLEX, 0.6, (0, 255,
92     0), 2)
93     else:
94         cv2.putText(output, 'computing...', (10, 25), cv2.FONT_HERSHEY_SIMPLEX, 0.6,
95         (0, 0, 255), 2)
96     # cv2.putText(output, text, (35, 50), cv2.FONT_HERSHEY_SIMPLEX,
97     # 1.25, (0, 255, 0), 2)
98
99     # check if the video writer is None
100    if writer is None:
101        # initialize our video writer
102        fourcc = cv2.VideoWriter_fourcc(*"MJPG")
103        writer = cv2.VideoWriter(args.output, fourcc, 30,
104        (W, H), True)
105
106    # write the output frame to disk
107    writer.write(output)
108
109    # show the output image
110    cv2.imshow("Output", output)
111    key = cv2.waitKey(1) & 0xFF
112
113    # if the 'q' key was pressed, break from the loop
114    if key == ord("q"):
115        break
116
117    # release the file pointers
118    print("[INFO] labeled video saved to the folder...")
119    writer.release()
120    vs.release()
121
122    # creating json file as: timeLabel.json
123    print("[INFO] writing json time-label...")
124    to_json_file(args.action, timeStamps)
125
126    print("[INFO] saving time-label figure to the folder...")
127    x_p = [item[0] for item in timeStamps]
128    y_p = [item[1] for item in timeStamps]
129    x_np_param = np.array(x_p)
130    y_np_param = np.array(y_p)
131
132    fig_plot(x_np_param, y_np_param, args.action)
133
134    print("[INFO] DONE...")
135
136
137
138 # python predict_video.py --model model/model_v5/activity.model --label-bin model/

```

```
model_v5/lb.pickle --input example_clips/v5/cook5.mp4 --output output/output_v5/
c5_v5.avi --action cooking --size 128
```

6 Running Instruction

6.1 Install Dependencies

- keras
- numpy
- opencv

7 Running the Prediction Code

```
1
2 python predict_video.py --model model/activity.model --label-bin model/lb.pickle
3 --input "YOUR VIDEO PATH" --output output/"ARBITRARY NAME".avi
4 --action drinking --size 128
```

7.1 Video Instruction Link

Playlist link >> [HERE](#) <<

or

https://www.youtube.com/watch?v=LGWF0j00oMI&list=PLHQn6Zmgbs7lScSfC6C0Ph_tF7Hu57teD0

8 Performance, Loss and Accuracy Figures

8.1 Previous Attempts

[INFO] evaluating network...

	precision	recall	f1-score	support
cooking	0.87	0.98	0.92	1958
drinking	0.95	0.71	0.81	977
accuracy			0.89	2935
macro avg	0.91	0.84	0.86	2935
weighted avg	0.90	0.89	0.88	2935

(a) Performance after 50 epochs

[INFO] evaluating network...

	precision	recall	f1-score	support
cooking	0.85	0.98	0.91	1958
drinking	0.94	0.66	0.77	977
accuracy			0.87	2935
macro avg	0.89	0.82	0.84	2935
weighted avg	0.88	0.87	0.87	2935

(b) Performance after 10 epochs

```
Epoch 00008: val_loss improved from 0.20288 to 0.19947, saving model to weights/weights.hdf5
Epoch 9/50
275/275 [=====] - 131s 478ms/step - loss: 0.2744 - accuracy: 0.8798
- val_loss: 0.3043 - val_accuracy: 0.8867

Epoch 00009: val_loss did not improve from 0.19947
Epoch 10/50
275/275 [=====] - 131s 476ms/step - loss: 0.2661 - accuracy: 0.8822
- val_loss: 0.0519 - val_accuracy: 0.8825

Epoch 00010: val_loss improved from 0.19947 to 0.05192, saving model to weights/weights.hdf5
Epoch 11/50
275/275 [=====] - 131s 476ms/step - loss: 0.2638 - accuracy: 0.8844
- val_loss: 0.1969 - val_accuracy: 0.8856

Epoch 00011: val_loss did not improve from 0.05192
Epoch 12/50
275/275 [=====] - 130s 475ms/step - loss: 0.2654 - accuracy: 0.8812
- val_loss: 0.05451 - val_accuracy: 0.8794

Epoch 00012: val_loss did not improve from 0.05192
Epoch 13/50
275/275 [=====] - 131s 475ms/step - loss: 0.2547 - accuracy: 0.8903
- val_loss: 0.3427 - val_accuracy: 0.8898
```

(c) Training procedure between epoch 8–12

```
Epoch 00045: val_loss did not improve from 0.05192
Epoch 46/50
275/275 [=====] - 130s 471ms/step - loss: 0.1999 - accuracy: 0.9116
- val_loss: 0.4214 - val_accuracy: 0.8753

Epoch 00046: val_loss did not improve from 0.05192
Epoch 47/50
275/275 [=====] - 131s 475ms/step - loss: 0.2029 - accuracy: 0.9144
- val_loss: 0.1752 - val_accuracy: 0.8915

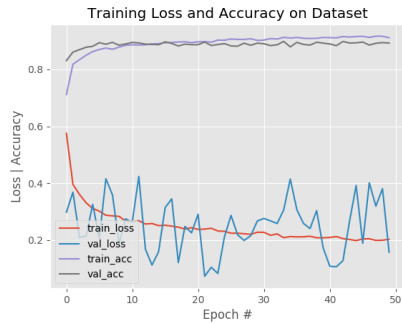
Epoch 00047: val_loss did not improve from 0.05192
Epoch 48/50
275/275 [=====] - 130s 473ms/step - loss: 0.2001 - accuracy: 0.9147
- val_loss: 0.7645 - val_accuracy: 0.8805

Epoch 00048: val_loss did not improve from 0.05192
Epoch 49/50
275/275 [=====] - 130s 474ms/step - loss: 0.1965 - accuracy: 0.9185
- val_loss: 0.2798 - val_accuracy: 0.8595

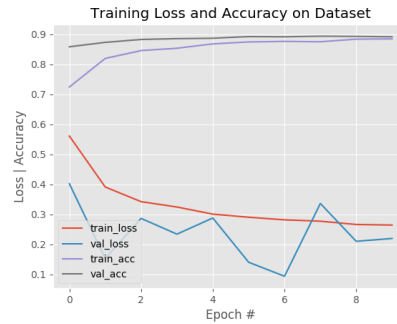
Epoch 00049: val_loss did not improve from 0.05192
Epoch 50/50
275/275 [=====] - 130s 474ms/step - loss: 0.2057 - accuracy: 0.9112
- val_loss: 0.2782 - val_accuracy: 0.8794

Epoch 00050: val_loss did not improve from 0.05192
```

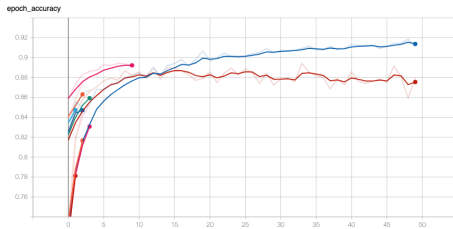
(d) Training procedure between epoch 45–50



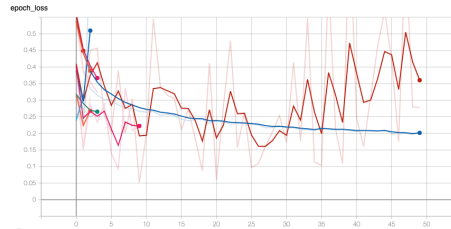
(e) Loss — Acc after 50 Epochs



(f) Loss–Acc after 10 Epochs



(g) Loss–Acc logs performance



(h) Epoch–Loss logs performance

8.2 Latest Improvements

```
[INFO] evaluating network...
```

	precision	recall	f1-score	support
cooking	0.87	0.94	0.90	1958
drinking	0.85	0.71	0.78	977
accuracy			0.86	2935
macro avg	0.86	0.83	0.84	2935
weighted avg	0.86	0.86	0.86	2935

(a) Performance-Final Project



(b) Loss-Acc-Final Project

```
Epoch 00030: val_loss did not improve from 0.44479
Epoch 31/50
275/275 [=====] - 131s 476ms/step - loss: 0.5372 - accuracy: 0.8398
- val_loss: 0.4174 - val_accuracy: 0.8753

Epoch 00031: val_loss improved from 0.44479 to 0.41739, saving model to weights/weights.hdf5
Epoch 32/50
275/275 [=====] - 133s 482ms/step - loss: 0.5356 - accuracy: 0.8425
- val_loss: 0.4753 - val_accuracy: 0.8553

Epoch 00032: val_loss did not improve from 0.41739
Epoch 33/50
275/275 [=====] - 132s 481ms/step - loss: 0.5325 - accuracy: 0.8390
- val_loss: 0.4421 - val_accuracy: 0.8622

Epoch 00033: val_loss did not improve from 0.41739
Epoch 34/50
275/275 [=====] - 128s 467ms/step - loss: 0.5309 - accuracy: 0.8371
- val_loss: 0.3705 - val_accuracy: 0.8650

Epoch 00034: val_loss improved from 0.41739 to 0.37049, saving model to weights/weights.hdf5
Epoch 35/50
275/275 [=====] - 131s 476ms/step - loss: 0.5289 - accuracy: 0.8362
- val_loss: 0.4811 - val_accuracy: 0.8643
```

(c) Training procedure between epoch 30-34-Final Project

```
Epoch 00040: val_loss did not improve from 0.37049
Epoch 41/50
275/275 [=====] - 128s 467ms/step - loss: 0.5073 - accuracy: 0.8533
- val_loss: 0.4072 - val_accuracy: 0.8674

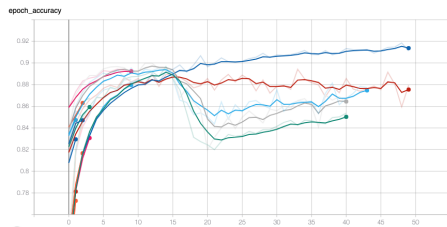
Epoch 00041: val_loss did not improve from 0.37049
Epoch 42/50
275/275 [=====] - 131s 476ms/step - loss: 0.5037 - accuracy: 0.8492
- val_loss: 0.5217 - val_accuracy: 0.8719

Epoch 00042: val_loss did not improve from 0.37049
Epoch 43/50
275/275 [=====] - 132s 479ms/step - loss: 0.5063 - accuracy: 0.8440
- val_loss: 0.4921 - val_accuracy: 0.8784

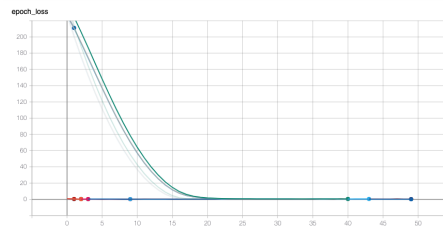
Epoch 00043: val_loss did not improve from 0.37049
Epoch 44/50
275/275 [=====] - 131s 475ms/step - loss: 0.5041 - accuracy: 0.8449
- val_loss: 0.5585 - val_accuracy: 0.8774

Epoch 00044: val_loss did not improve from 0.37049
```

(d) Training procedure between epoch 40-44-Final Project



(e) Loss-Acc logs performance-Final Project



(f) Epoch-Loss logs performance-Final Project